

FocusBuddy

PRD and Agile Backlog

1. Product Objective

The primary objective of FocusBuddy is to reduce the time between a user's intention to complete a task and their first meaningful action toward completing it. By lowering the cognitive friction associated with starting work, FocusBuddy helps users overcome task paralysis, build momentum, and make consistent progress toward their goals.

2. Problem and Goal

2.1 Problem Being Solved:

Traditional productivity systems require heavy cognitive overhead (sorting, assigning priorities, formatting lists) which actively triggers task freeze, demand avoidance, and executive overload in neurodivergent brains

2.2 User Goal:

To break out of painful, anxiety-inducing task paralysis loops instantly and initiate work consistently without relying on toxic, deadline-driven panic cycles.

2.3 Business Goal:

Achieve strong initial user traction and high 7-day organic retention within the neurodivergent and cognitive accessibility niche, establishing a baseline utility platform that can later scale into an enterprise-level B2B cognitive accessibility layer.

2.4 Product Goal:

Build a zero-onboarding, zero-setup conversational workflow that consumes unstructured, chaotic emotional input and returns an actionable, single-step plan within less than 2.5 seconds.

3. User Stories

User Story 1

Low-Friction Vent-Capture Engine:

As an overwhelmed university student, **I want to** voice-note or text my messy, unstructured thoughts into a non-judgmental interface, **so that I can** offload my mental panic without worrying about formatting a task perfectly.

User Story 2

Single-Step Progressive Interface:

As a neurodivergent remote worker, **I want** the system to hide the full scope of a massive project and only reveal one tiny step at a time, **so that I** don't experience secondary paralysis from looking at an intimidating backlog or checklist.

User Story 3

Dynamic Strategy Toggles:

As a user paralyzed by a boring or terrifying task, **I want to** flag my exact emotional block (e.g., "Low Energy," "Terrified," or "Avoiding"), **so that** the AI automatically adjusts its coaching persona, response tone, and initial step sizes to match my current cognitive capacity.

User Story 4

Dynamic Step Size Calibration:

As an anxious student, **I want** an immediate, friction-free way to signal that a suggested step still feels too difficult, **so that** the system dynamically down-scales the task into an even smaller, more absurdly easy micro-action.

4. Acceptance Criteria

User Story 1: Low-Friction Vent-Capture Engine

Scenario: Successful backend processing of a highly chaotic, unstructured user brain-dump.

Given the user is on the text capture screen and is experiencing high anxiety or task freeze

When the user inputs a chaotic, emotional block of text (e.g., "I have to do my taxes, but my room is messy and I can't find my documents, and everything sucks") and hits send

Then the system must process the text, strip out emotional filler phrases, isolate the primary core target item ("organize tax documents"), and transition the user to the single-step execution screen within 2.5 seconds.

User Story 2: Single-Step Progressive Interface

Scenario: Initializing the progressive disclosure flow while preventing cognitive overstimulation.

Given the Adaptive Micro-Step Engine has successfully parsed a target task into 5 logical, sequential micro-steps,

When the execution screen loads for the user,

Then the UI must display only Step 1 as a prominent, centered action item, while completely hiding Steps 2 through 5 from view to eliminate anticipatory overwhelm.

User Story 3: Dynamic Strategy Toggles

Scenario: User updates their current cognitive state via a strategy toggle to alter the AI's behavior.

Given the user is on the execution screen and has selected the "Terrified / Scared of Failing" toggle,

When the system generates or displays a micro-step,

Then the system must filter the step text through a low-pressure coaching prompt layout—stripping out any authoritative corporate tone or strict timelines and explicitly embedding a "permission to execute imperfectly" frame (e.g., "Open a scratchpad and write down three completely terrible ideas. Grammar doesn't count right now.").

User Story 4: Dynamic Step Size Calibration

Scenario: A heavily paralyzed user flags a generated micro-step as still being too overwhelming.

Given the system is currently displaying a standard micro-step: "Open your laptop and create a blank folder named Capstone Project.",

When the user clicks the "Too Hard / Can't Do It" calibration button.

Then the system must instantly intercept this signal, reduce the task friction threshold past the standard baseline, and display an alternate, micro-action below the "Absurdity Threshold": "Just physically go sit down at your desk. Don't touch your laptop yet. Let me know when you are seated."

5. Functional Requirements

1. Context Capture & Processing

- ✓ The system must accept unstructured raw text inputs up to 1,000 characters.
- ✓ The system must support a native microphone audio recording option for voice-note streaming wrappers, passing the file to an upstream text-to-speech engine.

2. AI Engine Interaction

- ✓ The system must parse user input using a secure system prompt that mandates a structured JSON output containing isolated, sequentially grouped micro-steps.
- ✓ The AI engine must apply a strict conversational persona that is objective, neutral, and supportive, explicitly stripping out toxic corporate urgency and patronizing fluff.

3. Execution Workspace UI

- ✓ The interface must display a singular task step at any given moment, masking all subsequent blocks.
- ✓ The UI must provide distinct, immediate action elements: [Done - Next Step], [Too Hard / Scale Down], and [Change Strategy Toggle].

6. Non-Functional Requirements

1. **Response Time (Performance):** The system must return the first parsed micro-step within less than or equal to 2.0 seconds of user submission to prevent attention from drifting and keep the interaction conversational.
2. **Accessibility (UI/UX):** The app interface must conform strictly to **WCAG 2.1 AA** guidelines, utilizing high-contrast styling, zero-clutter white space layouts, and clean screen-reader compatibility to minimize visual overstimulation.
3. **Privacy & Data Protection:** The product must employ a strict privacy-first data policy. User vent inputs, stress profiles, and cognitive states must be encrypted in transit and at rest, and explicitly excluded from upstream LLM training data loops via enterprise API agreements.
4. **Reliability & Fallbacks:** The core application wrapper must maintain a **99.9% uptime** standard, degrading gracefully to a localized, static set of standard grounding prompts if the underlying third-party LLM API experiences an outage or severe latency spike.

7. Assumptions

1. Users feel psychologically safe typing or speaking honestly to an AI entity about their active workflow anxieties, guilt, and emotional blockers.
2. Users possess access to a stable internet connection while actively trying to initiate their digital, professional, or academic tasks.
3. Users prefer an extremely stripped-back, near-blank execution window over a rich features dashboard during a moment of high cognitive crisis.
4. Users will not intentionally use the vent engine to generate or process dangerous, harmful, or out-of-scope for personal medical crises.

8. Dependencies

1. AI Services

- Availability of third-party LLM providers such as OpenAI or Anthropic.
- Consistent response quality and latency from AI APIs.

2. Design Resources

- UX design support for conversational flows.
- Accessibility review for cognitive-friendly interfaces.

3. Engineering Resources

- Frontend development for conversational interfaces.

- Backend infrastructure for AI orchestration and session management.

4. Authentication & User Management

- Support for email-based login, magic links, or social sign-in.
- Secure user session management.

5. Analytics & Monitoring

- Product analytics for activation and retention metrics.
- Monitoring tools for uptime, latency, and error tracking.

9. Risks & Mitigations

Risk 1: AI Generates Ineffective or Overwhelming Steps

Category: Product Risk

Description: The AI may generate steps that are still too large, vague, or difficult for users to act upon, reducing the product's effectiveness.

Mitigation: Implement prompt engineering guidelines, user feedback loops ("Too Hard"), and ongoing evaluation of generated micro-steps to improve output quality.

Risk 2: Low User Retention

Category: User Risk

Description: Users may initially find the product helpful but stop using it after the novelty wears off or if they do not perceive consistent value.

Mitigation: Measure activation success, introduce lightweight personalization, and continuously improve the quality and relevance of guidance.

Risk 3: AI Response Latency

Category: Technical Risk

Description: Slow AI responses may interrupt the conversational flow and reduce the product's usefulness during moments of overwhelm.

Mitigation: Use lightweight models for the MVP, optimize prompts, and implement fallback responses when latency exceeds acceptable thresholds.

Risk 4: Privacy and Trust Concerns

Category: User / Business Risk

Description: Users may hesitate to share personal thoughts, work challenges, or emotional concerns with an AI-powered system.

Mitigation: Clearly communicate privacy policies, minimize data collection, encrypt user data, and provide transparency around data usage.

Risk 5: Weak Problem-Solution Fit

Category: Product Risk

Description: The assumption that AI-generated micro-steps meaningfully increase task initiation may prove incorrect.

Mitigation: Track activation metrics, conduct user testing, and validate whether users start tasks more frequently after interacting with FocusBuddy.

10. Agile Backlog

Epic: Task Initiation Assistance

Enable users experiencing overwhelm or task paralysis to quickly transition from intention to action through AI-guided micro-steps.

Epic	User Story	Priority	Story Points	Notes
Task Initiation Assistance	As a user, I want to describe my task or concern in natural language so that I can receive help without organizing my thoughts first.	High	3	Supports free-form text input
Task Initiation Assistance	As a user, I want the system to identify the primary task from my input so that I can focus on execution instead of analysis.	High	5	AI task extraction capability
Task Initiation Assistance	As a user, I want to receive one manageable next step at a time so that I do not feel overwhelmed by large task lists.	High	5	Progressive disclosure workflow
Task Initiation Assistance	As a user, I want to indicate when a suggested step is still too difficult so that the system can generate a smaller action.	Medium	5	Adaptive micro-step refinement
Task Initiation Assistance	As a user, I want to mark a step as completed and receive the next action so	Medium	3	Continuous guidance loop

	that I can maintain momentum.			
--	-------------------------------	--	--	--